

Output of first step:

```
hafsa_027@Dell:~/k8s-master-app$ ./scripts/deploy.sh
=====
          KUBERNETES ZERO TO HERO - DEPLOYMENT
=====
[STEP 1] CHECKING PREREQUISITES
✓ minikube is installed
✓ kubectl is installed
✓ docker is installed
```

Step2:

```
# Step 2: Ensure Minikube is running
echo -e "${MAGENTA}[STEP 2] ENSURING MINIKUBE IS RUNNING${NC}"

if ! minikube status | grep -q "host: Running"; then
    echo -e "${YELLOW}Minikube is not running. Starting Minikube...${NC}"
    # Start Minikube with appropriate configuration
    minikube start --cpus=2 --memory=4096 --disk-size=20g

    if [ $? -ne 0 ]; then # Stores the exit status of the most recent command
        #0 → The exit code 0 typically means success.
        echo -e "${RED}Failed to start Minikube. Trying with minimal configuration...${NC}"
        # Fallback to minimal configuration
        minikube start --driver=docker #Uses Docker as the virtualization driver , Runs Kubernetes inside a Docker container instead of a virtual machine.

        if [ $? -ne 0 ]; then
            echo -e "${RED}Failed to start Minikube. Exiting.${NC}"
            exit 1
        fi
    else
        echo -e "${GREEN}✓ Minikube is already running${NC}"
    fi
fi
```

Output of Step 2:

```
v DOCKER is installed
[STEP 2] ENSURING MINIKUBE IS RUNNING
✓ Minikube is already running
```

Step3:

```
# Step 3: Enable required Minikube addons
echo -e "${MAGENTA}[STEP 3] ENABLING MINIKUBE ADDONS${NC}"

# We'll handle addons more carefully, checking if they're already enabled
# and handling errors better

# Function to safely enable an addon
enable_addon() {
    local addon=$1
    local already_enabled=$(minikube addons list | grep $addon | grep -c "enabled")

    if [ $already_enabled -eq 1 ]; then
        echo -e "${GREEN}✓ $addon addon is already enabled${NC}"
        return 0
    fi

    echo -e "${CYAN}Enabling $addon addon...${NC}"
    minikube addons enable $addon

    if [ $? -ne 0 ]; then
        echo -e "${YELLOW}Warning: Failed to enable $addon addon. Continuing without it.${NC}"
        return 1
    else
        echo -e "${GREEN}✓ $addon addon enabled${NC}"
        return 0
    fi
}

# Try to enable addons but don't fail if they don't work
enable_addon "dashboard" || true

echo -e "${YELLOW}Note: Skipping Ingress and Metrics Server addons as they may not work in all environments${NC}"
echo -e "${YELLOW}The application will still work without these addons${NC}"
```

Output of Step 3:

```
[STEP 3] ENABLING MINIKUBE ADDONS
/ dashboard addon is already enabled
Note: Skipping Ingress and Metrics Server addons as they may not work in all environments
The application will still work without these addons
```

Step4:

```
# Step 4: Configure Docker to use Minikube's Docker daemon
echo -e "${MAGENTA}[STEP 4] CONFIGURING DOCKER TO USE MINIKUBE${NC}"
echo -e "${CYAN}This allows us to build images directly into Minikube's registry${NC}"

#minikube docker-env prints environment variables needed to redirect Docker commands to Minikube's internal Docker daemon.
eval $(minikube docker-env)
if [ $? -ne 0 ]; then
    echo -e "${RED}Failed to configure Docker to use Minikube. Exiting.${NC}"
    exit 1
fi
echo -e "${GREEN}✓ Docker configured to use Minikube's registry${NC}"
```

Output of Step 4:

```
[STEP 4] CONFIGURING DOCKER TO USE MINIKUBE
This allows us to build images directly into Minikube's registry
```

Step5:

```
# Step 5: Build the Docker image
echo -e "${MAGENTA}[STEP 5] BUILDING DOCKER IMAGE${NC}"

echo -e "${CYAN}Building k8s-master-app:latest image...${NC}"
cd ~/k8s-master-app/app

# Added retry mechanism for Docker build with better network settings
MAX_ATTEMPTS=3
BUILD_SUCCESS=false

for ATTEMPT in $(seq 1 $MAX_ATTEMPTS); do
    echo -e "${YELLOW}Build attempt $ATTEMPT of $MAX_ATTEMPTS...${NC}"

    # Use host network for better connectivity in WSL
    docker build --network=host -t k8s-master-app:latest .
    # Builds an image named k8s-master-app:latest from the current directory (.).
    # --network=host to avoid network-related issues.

    if [ $? -eq 0 ]; then
        BUILD_SUCCESS=true
        break
    else
        echo -e "${YELLOW}Build attempt $ATTEMPT failed. Waiting before retry...${NC}"
        sleep 5
    fi
done

if [ "$BUILD_SUCCESS" != "true" ]; then
    echo -e "${RED}Failed to build Docker image after $MAX_ATTEMPTS attempts. Exiting.${NC}"
    exit 1
fi

echo -e "${GREEN}✓ Docker image built successfully${NC}"
docker images | grep k8s-master-app
```

Output of Step 5:

```
[STEP 5] BUILDING DOCKER IMAGE
Building k8s-master-app:latest image...
Build attempt 1 of 3...
[+] Building 3.6s (15/15) FINISHED
--> [internal] load build definition from Dockerfile
--> transferring dockerfile: 1.62kB
--> [internal] load metadata for docker.io/library/python:3.9
--> [auth] library/python:pull token for registry-1.docker.io
--> [internal] load .dockerrcignore
--> => transferring context: 2B
--> [1/9] FROM docker.io/library/python:3.9@sha256:5ea663alc6ba266fdcac5949d1d2ea364ce30a2da92a3df95bb3c01437633ad9
--> [internal] load build context
--> => transferring context: 25.22kB
--> CACHED [2/9] WORKDIR /app
--> CACHED [3/9] RUN mkdir -p /data /config /logs && chmod 777 /data /config /logs
--> CACHED [4/9] COPY requirements.txt .
--> CACHED [5/9] RUN pip install --no-cache-dir --upgrade pip && pip install --no-cache-dir -r requirements.txt
--> [6/9] COPY app.py .
--> [7/9] RUN chmod +x app.py
--> [8/9] RUN echo '#!/bin/sh' > /healthcheck.sh && echo 'curl -s http://localhost:5000/api/health || exit 1' >> /healthcheck.sh &&
--> [9/9] RUN useradd -m appuser && chown -R appuser:appuser /app /data /config /logs
--> exporting to image
--> => exporting layers
--> => writing image sha256:3fd49816bef32c436402dc39118316a21c2f6fd05acb5d433ba49580c8e5b40e
--> => naming to docker.io/library/k8s-master-app:latest
/ Docker image built successfully
k8s-master-app          latest      3fd49816bef3  Less than a second ago  1.02GB
```

Step6:

```
# Step 6: Apply Kubernetes manifests
echo -e "${MAGENTA}[STEP 6] DEPLOYING TO KUBERNETES${NC}"
cd ~/k8s-master-app

echo -e "${CYAN}Creating namespace...${NC}"
kubectl apply -f k8s/base/namespace.yaml

echo -e "${CYAN}Creating ConfigMap and sample files...${NC}"
kubectl apply -f k8s/volumes/volumes.yaml

echo -e "${CYAN}Creating ConfigMap for application settings...${NC}"
kubectl apply -f k8s/config/configmap.yaml

echo -e "${CYAN}Creating Secret...${NC}"
kubectl apply -f k8s/config/secret.yaml

echo -e "${CYAN}Creating Deployment...${NC}"
kubectl apply -f k8s/base/deployment.yaml
#Defines how the application pods should run, their container image, replicas, and resource limits.

echo -e "${CYAN}Creating Service...${NC}"
kubectl apply -f k8s/networking/service.yaml
#Creates a Service to expose the application for external access.

echo -e "${CYAN}Creating HorizontalPodAutoscaler...${NC}"
kubectl apply -f k8s/monitoring/hpa.yaml || true
if [ $? -ne 0 ]; then
  echo -e "${YELLOW}Warning: Failed to create HorizontalPodAutoscaler. This is expected if metrics-server is not enabled.${NC}"
  echo -e "${YELLOW}The application will still work without auto-scaling.${NC}"
fi

echo -e "${GREEN}✓ All Kubernetes resources applied${NC}"
```

Output of Step 6:

```
[STEP 6] DEPLOYING TO KUBERNETES
Creating namespace...
namespace/k8s-demo unchanged
Creating ConfigMap and sample files...
configmap/app-files unchanged
Creating ConfigMap for application settings...
configmap/app-config unchanged
Creating Secret...
secret/app-secrets unchanged
Creating Deployment...
deployment.apps/k8s-master-app configured
Creating Service...
service/k8s-master-app unchanged
Creating HorizontalPodAutoscaler...
horizontalpodautoscaler.autoscaling/k8s-master-hpa configured
✓ All Kubernetes resources applied
```

Step7:

```
# Step 7: Wait for deployment to be ready
echo -e "${MAGENTA}[STEP 7] WAITING FOR DEPLOYMENT TO BE READY${NC}"
echo -e "${CYAN}This may take a minute or two...${NC}"

echo "Waiting for deployment to be ready..."
kubectl -n k8s-demo rollout status deployment/k8s-master-app --timeout=180s

if [ $? -ne 0 ]; then
    echo -e "${RED}Deployment failed to become ready within the timeout period.${NC}"
    echo -e "${YELLOW}Checking pod status...${NC}"

    kubectl -n k8s-demo get pods

    echo -e "${YELLOW}Checking pod logs...${NC}"
    POD=$(kubectl -n k8s-demo get pods -l app=k8s-master -o name | head -1)
    # -l (or --selector): Filters pods based on labels.
    # -o (or --output): Controls the format of the output.
    # name: Returns the pod names in the format pod/<pod-name> instead of a detailed table.
    # head -1: Extracts only the first line

    if [ ! -z "$POD" ]; then
        #!: Negates the condition,-z "$POD": Checks if POD is empty (i.e., "" or not set).
        |   kubectl -n k8s-demo logs $POD
    fi
else
    echo -e "${GREEN}✓ Deployment is ready${NC}"
fi
```

Output of Step 7:

```
[STEP 7] WAITING FOR DEPLOYMENT TO BE READY
This may take a minute or two...
Waiting for deployment to be ready...
deployment "k8s-master-app" successfully rolled out
✓ Deployment is ready
```

Step8:

```
$ deploy.sh ×

203
204 # Step 8: Set up port forwarding for easier access
205 echo -e "${MAGENTA}[STEP 8] SETTING UP PORT FORWARDING${NC}"
206 echo -e "${CYAN}This will make the application accessible on localhost${NC}"
207
208 # To avoid conflicts from multiple port-forward processes running at the same time.
209 # Checks if a kubectl port-forward process is already running for k8s-demo.
210 # If found, pgrep returns the process ID(s).
211 # If a process is found (> /dev/null means we discard output):
212 # pkill -f "kubectl.*port-forward.*k8s-demo"
213 # Kills any existing port-forwarding processes related to k8s-demo.
214 if pgrep -f "kubectl.*port-forward.*k8s-demo" > /dev/null; then
215     echo -e "${YELLOW}Port forwarding is already running. Stopping it...${NC}"
216     pkill -f "kubectl.*port-forward.*k8s-demo"
217 fi
218
219 # Start port forwarding in the background
220 #Forwards port 80 on the service k8s-master-app
221 # To port 8080 on localhost., & Runs the process in the background.
222 kubectl -n k8s-demo port-forward svc/k8s-master-app 8080:80 &
223 PORT_FORWARD_PID=$! #$: captures the process ID (PID) of the last background command,Stores it in PORT_FORWARD_PID for later checks.
224
225 # Give it a moment to start,Pauses execution for 2 seconds to allow port forwarding to stabilize.
226 sleep 2
227
228 # Check if port forwarding started successfully
229 #/dev/null is a special file in Linux that discards any data written to it."BLACK HOLE"
230 #ensures only the exit status is checked, not the output.
231 #The script only cares whether the process exists (exit code), not the output.
232 # If process exists: ps -p $PORT_FORWARD_PID returns exit code 0 (success).
233 # If process doesn't exist: It returns a non-zero exit code (failure).
234 #By redirecting output to /dev/null, the script can check the exit code silently
235 if ! ps -p $PORT_FORWARD_PID > /dev/null; then
236     echo -e "${RED}Failed to start port forwarding.${NC}"
237 else
238     echo -e "${GREEN}✓ Port forwarding started on port 8080${NC}"
239 fi
```

Output of Step 8:

If facing this error:

```
[Deployment] [Step 8] [INFO] [STEP 8] SETTING UP PORT FORWARDING
This will make the application accessible on localhost
Unable to listen on port 8080: Listeners failed to create with the following errors: [unable to create listener: Error listen tcp4 127.0.0.1:8080
0: bind: address already in use unable to create listener: Error listen tcp6 [::1]:8080: bind: address already in use]
error: unable to listen on any of the requested ports: [{8080 5000}]

Failed to start port forwarding.
```

netstat -tulnp | grep :8080 #to get process id command to check

what is running on 8080 port if there was earlier

sudo pkill -9 <PID> # to actually kill the process

```
[STEP 8] SETTING UP PORT FORWARDING
This will make the application accessible on localhost
Forwarding from 127.0.0.1:8080 -> 5000
Forwarding from [::1]:8080 -> 5000
```

Step9:

```
# Step 9: Display access information
echo -e "${MAGENTA}[STEP 9] DEPLOYMENT COMPLETE${NC}"
echo -e "${BLUE}===== ${NC}"
echo -e "${GREEN}Kubernetes Zero to Hero application has been deployed!${NC}"
echo -e "${BLUE}===== ${NC}"

echo -e "${YELLOW}Your application is accessible via multiple methods:${NC}"
echo ""
echo -e "${CYAN}1. Port Forwarding:${NC}"
echo "    URL: http://localhost:8080"
echo "    (This is running in the background with PID $PORT_FORWARD_PID)"
echo ""

# Get Minikube IP
MINIKUBE_IP=$(minikube ip)
echo -e "${CYAN}2. NodePort:${NC}"
echo "    URL: http://$MINIKUBE_IP:30080"
echo ""

echo -e "${CYAN}3. Minikube Service URL:${NC}"
echo "    Run: minikube service k8s-master-app -n k8s-demo"
echo ""
```

Output of Step 9:

```
[STEP 9] DEPLOYMENT COMPLETE
=====
Kubernetes Zero to Hero application has been deployed!
=====
Your application is accessible via multiple methods:

1. Port Forwarding:
URL: http://localhost:8080
(This is running in the background with PID 58472)

2. NodePort:
URL: http://192.168.49.2:30080

3. Minikube Service URL:
Run: minikube service k8s-master-app -n k8s-demo
```

===== USEFUL COMMANDS =====

output after the main.py:

Kubernetes Zero to Hero v1.0.0

A comprehensive Kubernetes demonstration application

Pod Information

Instance ID: ce6ea312
Hostname: k8s-master-app-869578d457-bp7jk
Environment: demo
Request count: 7
Platform: Linux-5.15.167.4-microsoft-standard-WSL2-x86_64-with-glibc2.36
Uptime: 21382.9 seconds

Resource Usage

| CPU Usage | Memory | Disk | Requests |
|-----------|--------|------|----------|
| 6.9% | 54.3% | 1.5% | 7 |

```
#Checks if it is readable (os.access(PATH, os.R_OK)).  
# Stores this information in a dictionary under the volumes key.  
  
        'config': {  
            'path': CONFIG_PATH,  
            'mounted': os.path.exists(CONFIG_PATH) and os.access(CONFIG_PATH, os.R_OK)  
        },  
        'logs': {  
            'path': LOG_PATH,  
            'mounted': os.path.exists(LOG_PATH) and os.access(LOG_PATH, os.R_OK)  
        }  
    },  
    'timestamp': datetime.datetime.now().isoformat()  
}  
}
```

◆ Example Output

- If `LOG_PATH` is set to `/var/log/myapp`, then:

```
python  
  
Clog_file = "/var/log/myapp/app.log"
```

- If `LOG_PATH` is not set, then:

```
python  
  
Clog_file = "/logs/app.log"
```

Kubernetes Zero to Hero v1.0.0

A comprehensive Kubernetes demonstration application

Pod Information

Instance ID: ce6ea312

Hostname: k8s-master-app-869578d457-bp7jk

Environment: demo

Request count: 4

Platform: Linux-5.15.167.4-microsoft-standard-WSL2-x86_64-with-glibc2.36

Uptime: 20162.1 seconds

Resource Usage

| CPU Usage | Memory |
|-----------|--------|
| 7.9% | 54.4% |

| Disk | Requests |
|------|----------|
| 1.5% | 4 |

Mounted Volumes

Data Volume

Path: /data

Status: Successfully mounted

| Files: | |
|-------------------|-----------------------|
| hello.txt | <button>View</button> |
| info.txt | <button>View</button> |
| sample-config.txt | <button>View</button> |

Config Volume

Path: /config

Status: Mounted but empty

Logs Volume

```
app > app.py
100     x():
101     d:
102     </style>
103     ad:
104     y>
105     <div class="container">
106         <h1>{{ app_name }} <span class="badge badge-primary">v{{ app_version }}</span>
107         <p>A comprehensive Kubernetes demonstration application</p>
108
109         <div class="info-box">
110             <h2>Pod Information</h2>
111             <p><strong>Instance ID:</strong> {{ instance_id }}</p>
112             <p><strong>Hostname:</strong> {{ system_info.hostname }}</p>
113             <p><strong>Environment:</strong> <span class="badge badge-success">{{ env }}</span></p>
114             <p><strong>Request count:</strong> {{ request_count }}</p>
115             <p><strong>Platform:</strong> {{ system_info.platform }}</p>
116             <p><strong>Uptime:</strong> {{ system_info.uptime }}</p>
117         </div>
118
119     </div>
120
121     <div class="info-box">
122         <h2>Resource Usage</h2>
123         <div class="metrics">
124             <div class="metric-card">
125                 <div class="metric-label">CPU Usage</div>
126                 <div class="metric-value">{{ resource_usage.cpu_percent }}%</div>
127             </div>
128             <div class="metric-card">
129                 <div class="metric-label">Memory</div>
130                 <div class="metric-value">{{ resource_usage.memory_percent }}%</div>
131             </div>
132             <div class="metric-card">
133                 <div class="metric-label">Disk</div>
134                 <div class="metric-value">{{ resource_usage.disk_usage }}</div>
135             </div>
136             <div class="metric-card">
137                 <div class="metric-label">Requests</div>
138                 <div class="metric-value">{{ metrics.requests }}</div>
139             </div>
140         </div>
141     </div>
142
143     <div class="container">
144         <div class="info-box">
145             <h2>Mounted Volumes</h2>
146
147             <h3>Data Volume</h3>
148             <p><strong>Path:</strong> {{ volumes.data.path }}</p>
149             <p><strong>Status:</strong>
150                 {% if volumes.data.status == 'mounted' %}
151                     <span class="success">Successfully mounted</span>
152                 {% elif volumes.data.status == 'empty' %}
153                     <span class="warning">Mounted but empty</span>
154                 {% else %}
155                     <span class="error">Error: {{ volumes.data.error }}</span>
156                 {% endif %}
157             </p>
158
159             {% if volumes.data.files %}
160                 <div class="file-list">
161                     <h4>Files:</h4>
162                     {% for file in volumes.data.files %}
163                         <div class="file-item">
164                             <span>{{ file }}</span>
165                             <a href="/view-file?path={{ volumes.data.path }}/{{ file }}" class="view-link">View</a>
166                         </div>
167                     {% endfor %}
168                 </div>
169             {% endif %}
170         </div>
171     </div>
172
173 
```

Config Volume

Path: /config

Status: Mounted but empty

Logs Volume

Path: /logs

Status: Successfully mounted

Files:

- init.log View
- app.log View
- container.log View

Actions

```

app > app.py
100   ):
290
291     iv class="container">
327       <div class="info-box">
353         <h3>Config Volume</h3>
354         <p><strong>Path:</strong> {{ volumes.config.path }}</p>
355         <p><strong>Status:</strong>
356           {% if volumes.config.status == 'mounted' %}
357             | <span class="success">Successfully mounted</span>
358           {% elif volumes.config.status == 'empty' %}
359             | <span class="warning">Mounted but empty</span>
360           {% else %}
361             | <span class="error">Error: {{ volumes.config.error }}</span>
362           {% endif %}
363
364
365         {% if volumes.config.files %}
366           <div class="file-list">
367             <h4>Files:</h4>
368             {% for file in volumes.config.files %}
369               <div class="file-item">
370                 <span>{{ file }}</span>
371                 | <a href="/view-file?path={{ volumes.config.path }}/{{ file }}>{{ file }}</a>
372               </div>
373             {% endfor %}
374           </div>
375         {% endif %}
376
378   <h3>Logs Volume</h3>
379   <p><strong>Path:</strong> {{ volumes.logs.path }}</p>
380   <p><strong>Status:</strong>
381     {% if volumes.logs.status == 'mounted' %}
382       | <span class="success">Successfully mounted</span>
383     {% elif volumes.logs.status == 'empty' %}
384       | <span class="warning">Mounted but empty</span>
385     {% else %}
386       | <span class="error">Error: {{ volumes.logs.error }}</span>
387     {% endif %}
388
389     {% if volumes.logs.files %}
390       <div class="file-list">
391         <h4>Files:</h4>
392         {% for file in volumes.logs.files %}
393           <div class="file-item">
394             <span>{{ file }}</span>
395             | <a href="/view-file?path={{ volumes.logs.path }}/{{ file }}>{{ file }}</a>
396           </div>
397         {% endfor %}
398       </div>
399     {% endif %}
400
401

```

Logs Volume

Path: /logs

Status: Successfully mounted

Files:

- init.log View
- app.log View
- container.log View

Actions

```

378   <h3>Logs Volume</h3>
379   <p><strong>Path:</strong> {{ volumes.logs.path }}</p>
380   <p><strong>Status:</strong>
381     {% if volumes.logs.status == 'mounted' %}
382       | <span class="success">Successfully mounted</span>
383     {% elif volumes.logs.status == 'empty' %}
384       | <span class="warning">Mounted but empty</span>
385     {% else %}
386       | <span class="error">Error: {{ volumes.logs.error }}</span>
387     {% endif %}
388
389     {% if volumes.logs.files %}
390       <div class="file-list">
391         <h4>Files:</h4>
392         {% for file in volumes.logs.files %}
393           <div class="file-item">
394             <span>{{ file }}</span>
395             | <a href="/view-file?path={{ volumes.logs.path }}/{{ file }}>{{ file }}</a>
396           </div>
397         {% endfor %}
398       </div>
399     {% endif %}
400
401

```

Actions

[Create a File](#) [API Info](#) [Health Check](#) [Metrics](#)

Environment Variables

APP_NAME: Kubernetes Zero to Hero

APP_VERSION: 1.0.0

ENVIRONMENT: demo

DATA_PATH: /data

CONFIG_PATH: /config

LOG_PATH: /logs

SECRET_KEY: dev-sec...

```

327   <div class="info-box">
488     | {% endif %}
489   </div>
490
491   <div class="info-box">
492     <h2>Actions</h2>
493     <div class="nav-links">
494       <a href="/create-file" class="nav-link">Create a File</a>
495       <a href="/api/info" class="nav-link">API Info</a>
496       <a href="/api/health" class="nav-link">Health Check</a>
497       <a href="/api/metrics" class="nav-link">Metrics</a>
498     </div>
499   </div>
500
501   <div class="info-box">
502     <h2>Environment Variables</h2>
503     <p><strong>APP_NAME:</strong> {{ app_name }}</p>
504     <p><strong>APP_VERSION:</strong> {{ app_version }}</p>
505     <p><strong>ENVIRONMENT:</strong> {{ environment }}</p>
506     <p><strong>DATA_PATH:</strong> {{ data_path }}</p>
507     <p><strong>CONFIG_PATH:</strong> {{ config_path }}</p>
508     <p><strong>LOG_PATH:</strong> {{ log_path }}</p>
509     <p><strong>SECRET_KEY:</strong> {{ secret_key|truncate(10, True, '...') }}</p>
510   </div>
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525

```

```

@app.route('/api/metrics')
def get_metrics():
    """API endpoint for application metrics - useful for monitoring systems"""
    # Get basic resource usage stats
    cpu_percent = psutil.cpu_percent()
    memory_info = psutil.virtual_memory()
    disk_info = psutil.disk_usage('/')

    # Collect all metrics
    all_metrics = {
        'system': {
            'cpu_percent': cpu_percent,
            'memory_used_percent': memory_info.percent,
            'memory_used_mb': memory_info.used / (1024 * 1024),
            'memory_total_mb': memory_info.total / (1024 * 1024),
            'disk_used_percent': disk_info.percent,
            'disk_used_gb': disk_info.used / (1024**3),
            'disk_total_gb': disk_info.total / (1024**3)
        },
        'application': {
            'uptime_seconds': time.time() - start_time,
            'total_requests': metrics['requests'],
            'data_reads': metrics['data_reads'],
            'data_writes': metrics['data_writes'],
            'errors': metrics['errors']
        },
        'instance': {
            'id': INSTANCE_ID,
            'hostname': socket.gethostname()
        },
        'timestamp': datetime.datetime.now().isoformat()
    }

```

The screenshot shows a browser window with the URL `127.0.0.1:8080/api/metrics`. The page displays a JSON object representing system and application metrics. The JSON structure is as follows:

```

{
  "application": {
    "data_reads": 11,
    "data_writes": 1,
    "errors": 0,
    "total_requests": 7,
    "uptime_seconds": 22072.66371488571
  },
  "instance": {
    "hostname": "k8s-master-app-869578d457-bp7jk",
    "id": "ce6ea312"
  },
  "system": {
    "cpu_percent": 6.4,
    "disk_total_gb": 1006.853931427002,
    "disk_used_gb": 14.417072296142578,
    "disk_used_percent": 1.5,
    "memory_total_mb": 3802.8671875,
    "memory_used_mb": 1794.75,
    "memory_used_percent": 54.3
  },
  "timestamp": "2025-03-05T11:54:37.822627"
}

```

```

@app.route('/api/info')
def api_info():
    """API endpoint returning application information"""
    return jsonify({
        'app_name': APP_NAME,
        'version': APP_VERSION,
        'environment': ENVIRONMENT,
        'instance_id': INSTANCE_ID,
        'hostname': socket.gethostname(),
        'request_count': request_count,
        'uptime_seconds': time.time() - start_time,
    })
    # Kubernetes volumes allow persistent storage across container restarts.
    # Volumes are mounted inside a container as directories.
    # This lets the application store logs, config files, and data that persist beyond container restarts.
    # Volumes are defined in your Kubernetes YAML files (deployment.yaml or volumes.yaml):
    'volumes': [
        {
            'data': {
                'path': DATA_PATH,
                'mounted': os.path.exists(DATA_PATH) and os.access(DATA_PATH, os.R_OK)
            },
            # Checks if it is readable (os.access(PATH, os.R_OK)).
            # Stores this information in a dictionary under the volumes key.
            'config': {
                'path': CONFIG_PATH,
                'mounted': os.path.exists(CONFIG_PATH) and os.access(CONFIG_PATH, os.R_OK)
            },
            'logs': [
                {
                    'path': LOG_PATH,
                    'mounted': os.path.exists(LOG_PATH) and os.access(LOG_PATH, os.R_OK)
                }
            ],
            'timestamp': datetime.datetime.now().isoformat()
        }
    ]
}

```

```

{
  "app_name": "Kubernetes Zero to Hero",
  "environment": "demo",
  "hostname": "k8s-master-app-869578d457-bp7jk",
  "instance_id": "ce6ea312",
  "request_count": 7,
  "timestamp": "2025-03-05T11:54:18.452301",
  "uptime_seconds": 22053.293340444565,
  "version": "1.0.0",
  "volumes": [
    {
      "config": {
        "mounted": true,
        "path": "/config"
      },
      "data": {
        "mounted": true,
        "path": "/data"
      },
      "logs": {
        "mounted": true,
        "path": "/logs"
      }
    }
  ]
}

```

```

@app.route('/api/health')
def health_check():
    """Health check endpoint for Kubernetes liveness and readiness probes"""
    # Check if we can access our mounted volumes
    data_ok = os.path.exists(DATA_PATH) and os.access(DATA_PATH, os.R_OK)
    config_ok = os.path.exists(CONFIG_PATH) and os.access(CONFIG_PATH, os.R_OK)
    logs_ok = os.path.exists(LOG_PATH) and os.access(LOG_PATH, os.W_OK)

    # For a real application, you might check database connections,
    # cache availability, etc.

    # Overall health status
    is_healthy = data_ok and config_ok and logs_ok

    # Log health check results
    logger.info(f"Health check: {'PASS' if is_healthy else 'FAIL'}")

    response = {
        'status': 'healthy' if is_healthy else 'unhealthy',
        'checks': {
            'data_volume': 'accessible' if data_ok else 'inaccessible',
            'config_volume': 'accessible' if config_ok else 'inaccessible',
            'logs_volume': 'writable' if logs_ok else 'not writable'
        },
        'timestamp': datetime.datetime.now().isoformat(),
        'hostname': socket.gethostname()
    }

    # Set the HTTP status code based on health
    status_code = 200 if is_healthy else 503

    return jsonify(response), status_code

```

127.0.0.1:8080/api/health

Pretty-print □

```
{
  "checks": {
    "config_volume": "accessible",
    "data_volume": "accessible",
    "logs_volume": "writable"
  },
  "hostname": "k8s-master-app-869578d457-bp7jk",
  "status": "healthy",
  "timestamp": "2025-03-05T11:53:56.963211"
}
```